

OBL1 - Operating Systems

By: John Ivar Eriksen

1 - The process abstraction

1.1

Briefly describe what happens when a process is started from a program on disk. A mode switch from kernel- to user-mode must happen. Explain why this is necessary.

The first thing that happens is that the operating system (OS) initialize a process control block (PCB) to represent this new process.

Next, the OS allocates memory for the new process, and the proceeds to load the program into the computer's memory (RAM).

The OS then allocates a user-level stack for user-level executions, and a kernel level stack for system calls, interrupts and processor exceptions.

The mode switch from kernel- to user mode must happen in order to isolate the process from having full kernel access. This has to do with permissions and privileges of processes. If a process was allowed to run in kernel mode, it would have access to take control of the processes used to control and allocate system resources such as memory, CPU time and execute privileged instructions. This would enable manipulation of hardware and reading of memory, which in turn could be used by either malicious, or buggy, software in some detrimental way. E.g. being able to read memory would enable the reading of sensitive information such as passwords and encryption keys, or taking up all system resources for itself.

1.2

The struct `task_struct` is defined in `include\linux\sched.h` header file, in line 738 as per kernel version `linux-6.5.2`, 2023-09-08 (subject to change, i.e. varies between kernel versions).

a) The field name from this struct that stores the process ID:

The fields named `pid` are responsible for storing process ID, which is what the acronym stands for (process **id**).

There are many entries related `pid` in the header file.

b) The field name from this struct that keeps track of accumulated virtual memory:

The fields named "mm" are responsible for the virtual memory. The name ("mm") is derived from the term "memory management".

C) Name two other fields found via `top`.

Looking at the display output of `top`, I decided on the fields `PR` and `NI`, i.e. 3rd and 4th column.

```
top - 15:39:53 up 23:13,  4 users,  load average: 0,20, 0,45, 0,49
Tasks: 331 total,  1 running, 330 sleeping,  0 stopped,  0 zombie
%Cpu(s):  3,2 us,  1,6 sy,  0,0 ni, 95,0 id,  0,1 wa,  0,0 hi,  0,0 si,  0,0 st
MiB Mem : 15860,4 total,   500,7 free, 11563,1 used,  4581,4 buff/cache
MiB Swap: 2048,0 total,  1330,0 free,   718,0 used.  4297,4 avail Mem

  PID USER   PR   NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 1323 root    20    0 1084128 272440 204992 S   4,3   1,7   14:22.42 Xorg
112983 mars    20    0  744304 107676  89216 S   4,3   0,7    0:00.13 spectacle
 97810 mars    20    0 1133,2g 273592  71168 S   3,3   1,7    5:54.11 Discord
 77191 mars    20    0    11,6g   2,3g  61392 S   3,0  15,1   15:16.84 java
105972 mars    20    0 2045612 149008 119148 S   3,0   0,9    0:19.58 konsole
 2021 mars    20    0 2050712 184620 107460 S   2,6   1,1    8:59.94 kwin_x11
 63125 mars    20    0    12,5g  640152 201840 S   2,0   3,9   27:43.20 firefox
 91209 mars    20    0 3606084 150308  28232 S   1,3   0,9    1:35.58 spotify
 2072 mars    20    0 4741628 351320 138104 S   1,0   2,2   14:08.06 plasmashell
 2337 mars    20    0 3480504 110924  22656 S   1,0   0,7   15:17.05 tresorit-daemon
   530 root    19   -1   49600  15516  14236 S   0,7   0,1    0:27.01 systemd-journal
 64125 mars    20    0    18,9g 209364  92176 S   0,7   1,3    5:12.46 WebExtensions
 69689 mars    20    0    19,3g 557292 100408 S   0,7   3,4    7:45.12 Isolated Web Co
110614 mars    20    0   17224   5888   3712 R   0,7   0,0    0:05.97 top
 1728 mars    20    0   20616  10880   7168 S   0,3   0,1    0:00.72 systemd
 1743 mars     9  -11   89908  37520   8208 S   0,3   0,2    0:44.29 pipewire-pulse
 2018 mars    20    0   738472 94412  76912 S   0,3   0,6    0:05.31 kmsserver
 2055 mars    20    0   238148  23424  19840 S   0,3   0,1    0:06.99 kglobalaccel5
 2106 mars    20    0   713204  37764  30720 S   0,3   0,2    0:13.99 org_kde-powerde
 2159 mars    20    0  6164892 431656  74284 S   0,3   2,7    5:20.79 jetbrains-toolb
 2593 mars    20    0   699548  12780  10604 S   0,3   0,1    0:01.84 xdg-desktop-por
 63443 mars    20    0   158136  18304  16768 S   0,3   0,1    0:31.90 kio_http_cache_
 64753 mars    20    0  2916584 272416  90616 S   0,3   1,7    1:39.21 Isolated Web Co
 69848 mars    20    0 3201928 675324  93252 S   0,3   4,2    6:53.81 Isolated Web Co
```

1. `PR` (Priority)

- `PR` refers to "priority", and is found under the field name `static_prio`.
- `static_prio` represents the initial priority assigned to the process by the kernel, and is the priority "in the moment" from the point of view of the task scheduler.

2. `NI` (Nice Value)

- `NI` refers to "nice value", and is found under the field name `normal_prio`.
- This value acts as a suggestion to the kernel of what priority the process "should" have, and acts to influence the priority given by the kernel itself.
- The priority of a given process is usually illustrated, and simplified, as the sum `PR = 20 + NI`. The lower the number, the higher the priority.
- However, the kernel can change the PR value regardless of NI value (but not the NI value itself) if needed.

2 - Process memory and segments

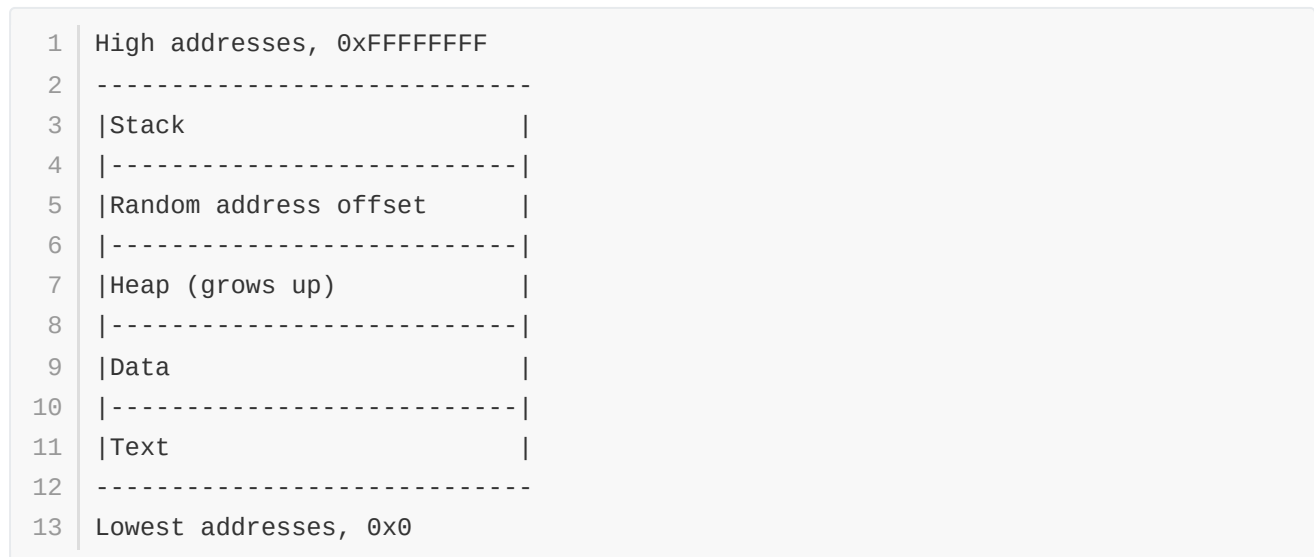
The memory region allocated to a process contains the following segments.

- Text segment
- Data segment
- Stack
- Heap

2.1

1. Sketch the organisation of a process' address space. Start with high addresses at the top, and the lowest address (0x0) at the bottom.

The address space of a process is organized as follows:



The layout contains some more segments, but these are the highlights. There's also some random address offsets for the Stack and Heap to ensure that they aren't found in easily predictable locations.

<https://manybutfinite.com/post/anatomy-of-a-program-in-memory/>

2.2

2. Briefly describe the purpose of each segment. Why is address 0x0 unavailable to the process?

Text segment

- Store binary image of the process, i.e. the executable code.

Data segment

- Global and static variables initialized by the program.

Stack

- Contains data that we know are going to be used.
- Local variables.
- Function call management and information. E.g. return addresses.

Heap

- Contains data that may be dynamically allocated while the program is running.
- Variables that may or may not be used.

The address 0x0 is unavailable to the process because it's reserved to act as a **null pointer**. E.g. when a process is trying to request memory, say 1 MB, but there's not enough available, the program will be sent the 0x0 address instead. This in turn is interpreted as "null pointer", and instead of causing a crash or other critical errors, the program simply receives this null pointer and can then throw a more beginning error message.

2.3.1

3. What are the differences between a global, static, and local variable?

Global variable

- Declared outside of any other functions/methods.
- A global "scope", i.e. not limited to any particular function, and accessible and modifiable by any functions.
- Automatically initialized to `0` when they are declared.

Static variable

- Scope is limited to the function/method, or block, in which it is declared.
- Retains its value between function calls.
- Only initialized once.
- Automatically initialized to `0`, unless initialized to something else.

Local variable

- Scope is limited to the method/function, or block, where it is defined.
- Available only inside the functions in which they are defined.
- Only exists until the function is executed. Once finished, the variable(s) are destroyed.

2.3.2

Given the following code snippet, show which segment each of the variables (var1, var2, var3) belong to.

```
1  #include<stdio.h>
2  #include <stdlib.h>
3
4  int var1 = 0;
5  void main()
6  {
7      int var2 = 1;
8      int *var3 = (int *)malloc(sizeof(int));
9      // Note, since we are using malloc(), var3 will be a
10     // pointer into the heap!
11     // So the question is, where is the pointer stored?
12     *var3 = 2;
13     printf("Address: %x; Value: %d\n", &var1, var1);
14     printf("Address: %x; Value: %d\n", &var2, var2);
```

```

15     printf("Address: %x; Address: %x; Value: %d\n", &var3, var3, *var3);
16 }

```

Output:

```

1 Address: 19f63014; Value: 0
2 Address: c0a8a2fc; Value: 1
3 Address: c0a8a300; "Address : 1a9592a0; Value: 2

```

var1

- Global variable; declared outside any functions.
- Typically stored in `data` segment.

var2

- Local variable; declared inside the `main` function.
- Typically stored on the `stack` segment.

var3

- Dynamically allocated variable, using `malloc` (**m**emory **a**llocation).
- Stored on the `heap` segment.

3 - Program code

3.1

1. Compile the example given above using `gcc mem.c -o mem`. Determine the sizes of the text, data, and bss segments using the command-line tool `size`.

Output of `size mem`:

```

1 text    data    bss     dec     hex filename
2 1801     616      8      2425    979 mem

```

3.2

2. Find the start address of the program using `objdump -f mem`.

Output of `objdump -f mem`:

```

1 mem:      file format elf64-x86-64
2 architecture: i386:x86-64, flags 0x000000150:
3 HAS_SYMS, DYNAMIC, D_PAGED
4 start address 0x000000000000010a0

```

3.3

3. Disassemble the compiled program using `objdump -d mem`. Capture the output and find the name of the function at the start address. Do a web search to find out what this function does, and why it is useful.

Output of `objdump -d mem`:

```
1 mem:      file format elf64-x86-64
2
3
4 Disassembly of section .init:
5
6 00000000000001000 <_init>:
7   1000:      f3 0f 1e fa      endbr64
8   1004:      48 83 ec 08      sub     $0x8,%rsp
9   1008:      48 8b 05 d9 2f 00 00  mov     0x2fd9(%rip),%rax      # 3fe8
   <__gmon_start__@Base>
10  100f:      48 85 c0          test    %rax,%rax
11  1012:      74 02            je      1016 <_init+0x16>
12  1014:      ff d0          call    *%rax
13  1016:      48 83 c4 08      add     $0x8,%rsp
14  101a:      c3              ret
15
16 Disassembly of section .plt:
17
18 00000000000001020 <.plt>:
19  1020:      ff 35 8a 2f 00 00  push    0x2f8a(%rip)      # 3fb0
   <_GLOBAL_OFFSET_TABLE_+0x8>
20  1026:      ff 25 8c 2f 00 00  jmp     *0x2f8c(%rip)      # 3fb8
   <_GLOBAL_OFFSET_TABLE_+0x10>
21  102c:      0f 1f 40 00      nopl    0x0(%rax)
22  1030:      f3 0f 1e fa      endbr64
23  1034:      68 00 00 00 00  push    $0x0
24  1039:      e9 e2 ff ff ff  jmp     1020 <_init+0x20>
25  103e:      66 90          xchg    %ax,%ax
26  1040:      f3 0f 1e fa      endbr64
27  1044:      68 01 00 00 00  push    $0x1
28  1049:      e9 d2 ff ff ff  jmp     1020 <_init+0x20>
29  104e:      66 90          xchg    %ax,%ax
30  1050:      f3 0f 1e fa      endbr64
31  1054:      68 02 00 00 00  push    $0x2
32  1059:      e9 c2 ff ff ff  jmp     1020 <_init+0x20>
33  105e:      66 90          xchg    %ax,%ax
34
35 Disassembly of section .plt.got:
36
37 00000000000001060 <__cxa_finalize@plt>:
38  1060:      f3 0f 1e fa      endbr64
39  1064:      ff 25 8e 2f 00 00  jmp     *0x2f8e(%rip)      # 3ff8
   <__cxa_finalize@GLIBC_2.2.5>
```

```

40      106a:      66 0f 1f 44 00 00      nopw    0x0(%rax,%rax,1)
41
42  Disassembly of section .plt.sec:
43
44  00000000000001070 <__stack_chk_fail@plt>:
45      1070:      f3 0f 1e fa      endbr64
46      1074:      ff 25 46 2f 00 00      jmp     *0x2f46(%rip)          # 3fc0
<__stack_chk_fail@GLIBC_2.4>
47      107a:      66 0f 1f 44 00 00      nopw    0x0(%rax,%rax,1)
48
49  00000000000001080 <printf@plt>:
50      1080:      f3 0f 1e fa      endbr64
51      1084:      ff 25 3e 2f 00 00      jmp     *0x2f3e(%rip)          # 3fc8
<printf@GLIBC_2.2.5>
52      108a:      66 0f 1f 44 00 00      nopw    0x0(%rax,%rax,1)
53
54  00000000000001090 <malloc@plt>:
55      1090:      f3 0f 1e fa      endbr64
56      1094:      ff 25 36 2f 00 00      jmp     *0x2f36(%rip)          # 3fd0
<malloc@GLIBC_2.2.5>
57      109a:      66 0f 1f 44 00 00      nopw    0x0(%rax,%rax,1)
58
59  Disassembly of section .text:
60
61  000000000000010a0 <_start>:
62      10a0:      f3 0f 1e fa      endbr64
63      10a4:      31 ed      xor     %ebp,%ebp
64      10a6:      49 89 d1      mov     %rdx,%r9
65      10a9:      5e      pop     %rsi
66      10aa:      48 89 e2      mov     %rsp,%rdx
67      10ad:      48 83 e4 f0      and     $0xfffffffffffffff0,%rsp
68      10b1:      50      push    %rax
69      10b2:      54      push    %rsp
70      10b3:      45 31 c0      xor     %r8d,%r8d
71      10b6:      31 c9      xor     %ecx,%ecx
72      10b8:      48 8d 3d ca 00 00 00      lea     0xca(%rip),%rdi          # 1189
<main>
73      10bf:      ff 15 13 2f 00 00      call    *0x2f13(%rip)          # 3fd8
<__libc_start_main@GLIBC_2.34>
74      10c5:      f4      hlt
75      10c6:      66 2e 0f 1f 84 00 00      cs nopw 0x0(%rax,%rax,1)
76      10cd:      00 00 00
77
78  000000000000010d0 <deregister_tm_clones>:
79      10d0:      48 8d 3d 39 2f 00 00      lea     0x2f39(%rip),%rdi          # 4010
<__TMC_END__>
80      10d7:      48 8d 05 32 2f 00 00      lea     0x2f32(%rip),%rax          # 4010
<__TMC_END__>
81      10de:      48 39 f8      cmp     %rdi,%rax
82      10e1:      74 15      je      10f8
<deregister_tm_clones+0x28>

```

```

83      10e3:      48 8b 05 f6 2e 00 00      mov     0x2ef6(%rip),%rax      # 3fe0
      <_ITM_deregisterTMCloneTable@Base>
84      10ea:      48 85 c0                    test    %rax,%rax
85      10ed:      74 09                        je      10f8
      <deregister_tm_clones+0x28>
86      10ef:      ff e0                        jmp     *%rax
87      10f1:      0f 1f 80 00 00 00 00      nopl    0x0(%rax)
88      10f8:      c3                          ret
89      10f9:      0f 1f 80 00 00 00 00      nopl    0x0(%rax)
90
91      00000000000001100 <register_tm_clones>:
92      1100:      48 8d 3d 09 2f 00 00      lea     0x2f09(%rip),%rdi      # 4010
      <__TMC_END__>
93      1107:      48 8d 35 02 2f 00 00      lea     0x2f02(%rip),%rsi      # 4010
      <__TMC_END__>
94      110e:      48 29 fe                    sub     %rdi,%rsi
95      1111:      48 89 f0                    mov     %rsi,%rax
96      1114:      48 c1 ee 3f                shr     $0x3f,%rsi
97      1118:      48 c1 f8 03                sar     $0x3,%rax
98      111c:      48 01 c6                    add     %rax,%rsi
99      111f:      48 d1 fe                    sar     %rsi
100     1122:      74 14                        je      1138 <register_tm_clones+0x38>
101     1124:      48 8b 05 c5 2e 00 00      mov     0x2ec5(%rip),%rax      # 3ff0
      <_ITM_registerTMCloneTable@Base>
102     112b:      48 85 c0                    test    %rax,%rax
103     112e:      74 08                        je      1138 <register_tm_clones+0x38>
104     1130:      ff e0                        jmp     *%rax
105     1132:      66 0f 1f 44 00 00          nopw    0x0(%rax,%rax,1)
106     1138:      c3                          ret
107     1139:      0f 1f 80 00 00 00 00      nopl    0x0(%rax)
108
109     00000000000001140 <__do_global_dtors_aux>:
110     1140:      f3 0f 1e fa                endbr64
111     1144:      80 3d c5 2e 00 00 00      cmpb    $0x0,0x2ec5(%rip)      # 4010
      <__TMC_END__>
112     114b:      75 2b                        jne     1178
      <__do_global_dtors_aux+0x38>
113     114d:      55                          push    %rbp
114     114e:      48 83 3d a2 2e 00 00      cmpq    $0x0,0x2ea2(%rip)      # 3ff8
      <__cxa_finalize@GLIBC_2.2.5>
115     1155:      00
116     1156:      48 89 e5                    mov     %rsp,%rbp
117     1159:      74 0c                        je      1167
      <__do_global_dtors_aux+0x27>
118     115b:      48 8b 3d a6 2e 00 00      mov     0x2ea6(%rip),%rdi      # 4008
      <__dso_handle>
119     1162:      e8 f9 fe ff ff            call    1060 <__cxa_finalize@plt>
120     1167:      e8 64 ff ff ff            call    10d0 <deregister_tm_clones>
121     116c:      c6 05 9d 2e 00 00 01      movb    $0x1,0x2e9d(%rip)      # 4010
      <__TMC_END__>
122     1173:      5d                          pop     %rbp
123     1174:      c3                          ret

```



```

124      1175:      0f 1f 00      nopl    (%rax)
125      1178:      c3            ret
126      1179:      0f 1f 80 00 00 00 00      nopl    0x0(%rax)
127
128 00000000000001180 <frame_dummy>:
129      1180:      f3 0f 1e fa      endbr64
130      1184:      e9 77 ff ff ff      jmp     1100 <register_tm_clones>
131
132 00000000000001189 <main>:
133      1189:      f3 0f 1e fa      endbr64
134      118d:      55              push    %rbp
135      118e:      48 89 e5          mov     %rsp,%rbp
136      1191:      48 83 ec 20        sub     $0x20,%rsp
137      1195:      64 48 8b 04 25 28 00      mov     %fs:0x28,%rax
138      119c:      00 00
139      119e:      48 89 45 f8        mov     %rax,-0x8(%rbp)
140      11a2:      31 c0             xor     %eax,%eax
141      11a4:      c7 45 ec 01 00 00 00      movl    $0x1,-0x14(%rbp)
142      11ab:      bf 04 00 00 00      mov     $0x4,%edi
143      11b0:      e8 db fe ff ff      call    1090 <malloc@plt>
144      11b5:      48 89 45 f0        mov     %rax,-0x10(%rbp)
145      11b9:      48 8b 45 f0        mov     -0x10(%rbp),%rax
146      11bd:      c7 00 02 00 00 00      movl    $0x2,(%rax)
147      11c3:      8b 05 4b 2e 00 00      mov     0x2e4b(%rip),%eax      # 4014
    <var1>
148      11c9:      89 c2            mov     %eax,%edx
149      11cb:      48 8d 05 42 2e 00 00      lea     0x2e42(%rip),%rax      # 4014
    <var1>
150      11d2:      48 89 c6          mov     %rax,%rsi
151      11d5:      48 8d 05 2c 0e 00 00      lea     0xe2c(%rip),%rax      # 2008
    <_IO_stdin_used+0x8>
152      11dc:      48 89 c7          mov     %rax,%rdi
153      11df:      b8 00 00 00 00      mov     $0x0,%eax
154      11e4:      e8 97 fe ff ff      call    1080 <printf@plt>
155      11e9:      8b 55 ec          mov     -0x14(%rbp),%edx
156      11ec:      48 8d 45 ec          lea     -0x14(%rbp),%rax
157      11f0:      48 89 c6          mov     %rax,%rsi
158      11f3:      48 8d 05 0e 0e 00 00      lea     0xe0e(%rip),%rax      # 2008
    <_IO_stdin_used+0x8>
159      11fa:      48 89 c7          mov     %rax,%rdi
160      11fd:      b8 00 00 00 00      mov     $0x0,%eax
161      1202:      e8 79 fe ff ff      call    1080 <printf@plt>
162      1207:      48 8b 45 f0        mov     -0x10(%rbp),%rax
163      120b:      8b 08             mov     (%rax),%ecx
164      120d:      48 8b 55 f0        mov     -0x10(%rbp),%rdx
165      1211:      48 8d 45 f0        lea     -0x10(%rbp),%rax
166      1215:      48 89 c6          mov     %rax,%rsi
167      1218:      48 8d 05 01 0e 00 00      lea     0xe01(%rip),%rax      # 2020
    <_IO_stdin_used+0x20>
168      121f:      48 89 c7          mov     %rax,%rdi
169      1222:      b8 00 00 00 00      mov     $0x0,%eax
170      1227:      e8 54 fe ff ff      call    1080 <printf@plt>

```

```

171      122c:      90                nop
172      122d:      48 8b 45 f8       mov     -0x8(%rbp),%rax
173      1231:      64 48 2b 04 25 28 00 sub     %fs:0x28,%rax
174      1238:      00 00
175      123a:      74 05            je      1241 <main+0xb8>
176      123c:      e8 2f fe ff ff    call    1070 <__stack_chk_fail@plt>
177      1241:      c9              leave   %eax
178      1242:      c3              ret
179
180  Disassembly of section .fini:
181
182  00000000000001244 <_fini>:
183      1244:      f3 0f 1e fa      endbr64
184      1248:      48 83 ec 08       sub     $0x8,%rsp
185      124c:      48 83 c4 08       add     $0x8,%rsp
186      1250:      c3              ret

```

The function at the start address is `main`, located at `0x1189` (`00000000000001189`). This is the entry point of the program, i.e. where execution of the program code begins.

The "lower" addresses are part of the programs initialization and setup, i.e. preparing the system for running the program.

3.4

4. Run the program several times (hint: running a program from the current directory is done using the syntax `./mem`). The addresses change between consecutive runs. Why?

The addresses change between consecutive runs due to randomization of the address space. This is a security feature to prevent e.g. malicious software from predicting where any given process will store its data in memory. This feature is called [Address Space Layout Randomization](#), or ASLR for short.

4 - The stack

Consider the following C program:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  // Changed pointers from 0x%08x to %p due to 64bit system.
4  void func()
5  {
6      char b = 'b';
7      /*long localvar = 2;
8      printf("func() with localvar @ %p\n", (void*)&localvar);
9      printf("func() frame address @ %p\n", (void*)__builtin_frame_address(0));
10     localvar++;*/
11     b = 'a';
12     func();
13 }
14

```

```

15 int main()
16 {
17     printf("main() frame address @ %p\n", __builtin_frame_address(0));
18     func();
19     exit(0);
20 }

```

4.1

1. Compile the example given above using `gcc stackoverflow.c -o stackoverflow`.

I had to change the pointers from `0x%08x` to `%p` due to running Linux on a 64 bit system, not in a 32 bit VM. A comment was made in the code pasted above.

4.2

2. Determine the default size of the stack for your Linux system. Hint: use the `ulimit` command (a web search or running the command `ulimit --help` will help find the appropriate command-line flags).

Output of `ulimit -s`:

```

1 8192

```

The output show the size of the stack in kilobytes.

4.3

3. Run the program. Describe your observations and find the cause of the error.

Output from running the `./stackoverflow` C-file:

```

1 main() frame address @ 0x7ffd5689fe00
2 Segmentation fault (core dumped)

```

The segmentation fault, or `segfault`, is a fault condition thrown by the memory protection system to notify the OS that the program has attempted to access a restricted memory area ("memory access violation"). This is kind of fault is often the result of programming mistakes that result in stack overflow error.

The reason for this is usually that the program is stuck in an infinite loop, where function calls or writing are running without any conditions under which it will terminate. It will continue using up stack space until it runs out, and since it's still calling functions it will try to write access memory outside the bounds of stack. This violates the restrictions on memory access, and is terminated with a `segfault`.

The reason for this, in this program, is the recursive calls on `func()`. Since there's no conditions set in `func()` where it will self-terminate, or other wise finish or end, the recursive funtion calls will run in an infinite loop.

4.4

4. Run the program and pipe the output to grep and wc -l:

```
./stackoverflow | grep func | wc -l
```

What does this number tell you about the stack? How does this relate to the default stack size you found using the ulimit command?

Output from running the program as is resultet in the output 0.

Assuming there is supposed to be something more to see, uncommenting line 9,

```
1 | printf("func() frame address @ %p\n", (void*)__builtin_frame_address(0));
```

results in the output 261712.

Further uncommenting line 6, 8 and 11:

```
1 | long localvar = 2;  
2 | printf("func() with localvar @ %p\n", (void*)&localvar);  
3 | localvar++;
```

outputs 349022.

The `wc -l` command outputs "word count, lines". If we focus on the second output, this means that a line with the string `func` was output 261 712 times before the segfault occurred. In short, the stack could fit 261 712 lines of the string `func() frame address @ %p\n` into the stack.

4.5

5. How much stack memory (in bytes) does each recursive function call occupy?

- Stack size (binary): 8 192 KiB * 1024 = 8 388 608 bytes
- Stack size (SI): 8192 kB * 1000 = 8 192 000 bytes
- Lines output: 261 712

Calculating stack frame size using binary size:

$$\frac{\text{Total stack size, bytes}}{\text{Number of calls}} = \frac{8\,388\,608}{261\,712} = 32.053 \text{ byte}$$

The calculation show that each recursive call occupy about 32 bytes of stack space.